

2. Basic Tricks	4
2.1. Saving State.....	4
2.2. Add, Delete, Rename	4
2.3. Advanced Undo/Redo	5
2.4. Reverting	6
2.5. Changelog Generation.....	6
2.6. Downloading Files	6
2.7. The Bleeding Edge.....	7
2.8. Instant Publishing.....	7
2.9. What Have I Done?.....	8
2.10. Exercise.....	8
3. Cloning Around	10
3.1. Sync Computers	10
3.2. Classic Source Control.....	10
3.3. Secret Source.....	11
3.4. Bare repositories	12
3.5. Push versus pull.....	12
3.6. Forking a Project.....	12
3.7. Ultimate Backups.....	13
3.8. Light-Speed Multitask.....	13
3.12. Why I use Git	15
4. Branch Wizardry.....	17
4.1. The Boss Key	17
4.2. Dirty Work	18
4.3. Quick Fixes	18
4.4. Merging.....	19
4.5. Uninterrupted Workflow	20
4.6. Reorganizing a Medley	21
4.7. Managing Branches.....	21
4.8. Temporary Branches	22
4.9. Work How You Want	22
iii5. Lessons of History.....	23
5.1. I Stand Corrected	23
5.2. . . . And Then Some.....	23
5.3. Local Changes Last.....	24
5.4. Rewriting History.....	25
5.5. Making History	25
5.6. Where Did It All Go Wrong?.....	27
5.7. Who Made It All Go Wrong?	27
5.8. Personal Experience.....	28
6. Multiplayer Git	29
6.1. Who Am I?.....	29
6.2. Git Over SSH, HTTP	29
6.3. Git Over Anything.....	30
6.4. Patches: The Global Currency	30
6.5. Sorry, We've Moved.....	31
6.6. Remote Branches	32
6.7. Multiple Remotes.....	33
6.8. My Preferences	33
7. Git Grandmastery.....	35
7.1. Source Releases.....	35
7.2. Commit What Changed.....	35
7.3. My Commit Is Too Big!.....	35
7.4. The Index: Git's Staging Area	36

7.5. Don't Lose Your HEAD.....	36
7.6. HEAD-hunting	37
7.7. Building On Git.....	38
7.8. Daring Stunts.....	39
7.9. Preventing Bad Commits	39

Chapter 2. Basic Tricks

Rather than diving into a sea of Git commands, use these elementary examples to get your feet wet. Despite their simplicity, each of them are useful. Indeed, in my first months with Git I never ventured beyond the material in this chapter.

2.1. Saving State

About to attempt something drastic? Before you do, take a snapshot of all files in the current directory with:

```
$ git init
$ git add .
$ git commit -m "My first backup"
```

Now if your new edits go awry, restore the pristine version:

```
$ git reset --hard
To save the state again:
$ git commit -a -m "Another backup"
```

2.2. Add, Delete, Rename

The above only keeps track of the files that were present when you first ran `git add`. If you add new files or subdirectories, you'll have to tell Git:

```
$ git add readme.txt Documentation
```

Similarly, if you want Git to forget about certain files:

```
$ git rm kludge.h obsolete.c
$ git rm -r incriminating/evidence/
```

Git deletes these files for you if you haven't already.

Renaming a file is the same as removing the old name and adding the new name. There's also the shortcut `git mv` which has the same syntax as the `mv` command. For example:

```
$ git mv bug.c feature.c
```

2.3. Advanced Undo/Redo

Sometimes you just want to go back and forget about every change past a certain point because they're all wrong.

Then:

```
$ git log
```

shows you a list of recent commits, and their SHA1 hashes:

```
commit 766f9881690d240ba334153047649b8b8f11c664
```

```
Author: Bob <bob@example.com>
```

```
Date: Tue Mar 14 01:59:26 2000 -0800
```

```
Replace printf() with write().
```

```
commit 82f5ea346a2e651544956a8653c0f58dc151275c
```

```
Author: Alice <alice@example.com>
```

```
Date: Thu Jan 1 00:00:00 1970 +0000
```

```
Initial commit.
```

The first few characters of the hash are enough to specify the commit; alternatively, copy and paste the entire hash. Type:

```
$ git reset --hard 766f
```

to restore the state to a given commit and erase all newer commits from the record permanently.

Other times you want to hop to an old state briefly. In this case, type:

```
$ git checkout 82f5
```

This takes you back in time, while preserving newer commits. However, like time travel in a science-fiction movie, if you now edit and commit, you will be in an alternate reality, because your actions are different to what they were the first time around.

This alternate reality is called a *branch*, and we'll have more to say about this later. For now, just remember that

```
$ git checkout master
```

will take you back to the present. Also, to stop Git complaining, always commit or reset your changes before running checkout.

To take the computer game analogy again:

- `git reset --hard`: load an old save and delete all saved games newer than the one just loaded.

- **git checkout:** load an old game, but if you play on, the game state will deviate from the newer saves you made the first time around. Any saved games you make now will end up in a separate branch representing the alternate reality you have entered. We deal with this later.

You can choose only to restore particular files and subdirectories by appending them after the command:

```
$ git checkout 82f5 some.file another.file
```

Take care, as this form of checkout can silently overwrite files. To prevent accidents, commit before running any checkout command, especially when first learning Git. In general, whenever you feel unsure about any operation, Git command or not, first run `git commit -a`.

Don't like cutting and pasting hashes? Then use:

```
$ git checkout :/"My first b"
```

to jump to the commit that starts with a given message. You can also ask for the 5th-last saved state:

```
$ git checkout master~5
```

2.4. Reverting

In a court of law, events can be stricken from the record. Likewise, you can pick specific commits to undo.

```
$ git commit -a
```

```
$ git revert 1b6d
```

will undo just the commit with the given hash. The revert is recorded as a new commit, which you can confirm by running `git log`.

2.5. Changelog Generation

Some projects require a changelog (<http://en.wikipedia.org/wiki/Changelog>). Generate one by typing:

```
$ git log > ChangeLog
```

2.6. Downloading Files

Get a copy of a project managed with Git by typing: `$ git`

```
clone git://server/path/to/files
```

For example, to get all the files I used to create this site: `$ git`

```
clone git://git.or.cz/gitmagic.git
```

We'll have much to say about the clone command soon.

2.7. The Bleeding Edge

If you've already downloaded a copy of a project using `git clone`, you can upgrade to the latest version with:

```
$ git pull
```

2.8. Instant Publishing

Suppose you've written a script you'd like to share with others. You could just tell them to download from your computer, but if they do so while you're improving the script or making experimental changes, they could wind up in trouble. Of course, this is why release cycles exist. Developers may work on a project frequently, but they only make the code available when they feel it is presentable.

To do this with Git, in the directory where your script resides:

```
$ git init
```

```
$ git add .
```

```
$ git commit -m "First release"
```

Then tell your users to run:

```
$ git clone your.computer:/path/to/script
```

to download your script. This assumes they have ssh access. If not, run `git daemon` and tell your users to instead run:

```
$ git clone git://your.computer/path/to/script
```

From now on, every time your script is ready for release, execute:

```
$ git commit -a -m "Next release"
```

and your users can upgrade their version by changing to the directory containing your script and typing:

```
$ git pull
```

Your users will never end up with a version of your script you don't want them to see.

2.9. What Have I Done?

Find out what changes you've made since the last commit with:

```
$ git diff
```

Or since yesterday:

```
$ git diff "@{yesterday}"
```

Or between a particular version and 2 versions ago:

```
$ git diff 1b6d "master~2"
```

In each case the output is a patch that can be applied with `git apply`. Try also:

```
$ git whatchanged --since="2 weeks ago"
```

Often I'll browse history with `qgit` (<http://sourceforge.net/projects/qgit>) instead, due to its slick photogenic interface, or `tig` (<http://jonas.nitro.dk/tig/>), a text-mode interface that works well over slow connections.

Alternatively, install a web server, run `git instaweb` and fire up any web browser.

Chapter 3. Cloning Around

In older version control systems, checkout is the standard operation to get files. You retrieve a bunch of files in a particular saved state.

In Git and other distributed version control systems, cloning is the standard operation. To get files, you create a *clone* of the entire repository. In other words, you practically mirror the central server. Anything the main repository can do, you can do.

3.1. Sync Computers

I can tolerate making tarballs or using `rsync` for backups and basic syncing. But sometimes I edit on my laptop, other times on my desktop, and the two may not have talked to each other in between.

Initialize a Git repository and commit your files on one machine. Then on the other:

```
$ git clone other.computer:/path/to/files
```

to create a second copy of the files and Git repository. From now on,

```
$ git commit -a
```

```
$ git pull other.computer:/path/to/files HEAD
```

will *pull* in the state of the files on the other computer into the one you're working on. If you've recently made conflicting edits in the same file, Git will let you know and you should commit again after resolving them.

3.2. Classic Source Control

Initialize a Git repository for your files:

```
$ git init
```

```
$ git add .
```

```
$ git commit -m "Initial commit"
```

On the central server, initialize a *bare repository* in some directory:

```
$ mkdir proj.git
```

```
$ cd proj.git
```

```
$ git --bare init
```

```
$ touch proj.git/git-daemon-export-ok
```

Start the Git daemon if necessary:

```
$ git daemon --detach # it may already be running
```

For Git hosting services, follow the instructions to setup the initially empty Git repository. Typically one fills in a form on a webpage.

Push your project to the central server with:

```
$ git push central.server/path/to/proj.git HEAD
```

To check out the source, a developer types:

```
$ git clone central.server/path/to/proj.git
```

After making changes, the developer saves changes locally:

```
$ git commit -a
```

To update to the latest version:

```
$ git pull
```

Any merge conflicts should be resolved then committed:

```
$ git commit -a
```

To check in local changes into the central repository:

```
$ git push
```

If the main server has new changes due to activity by other developers, the push fails, and the developer should pull the latest version, resolve any merge conflicts, then try again.

Developers must have SSH access for the above pull and push commands. However, anyone can see the source by typing:

```
$ git clone git://central.server/path/to/proj.git
```

The native git protocol is like HTTP: there is no authentication, so anyone can retrieve the project. Accordingly, by default, pushing is forbidden via the git protocol.

3.3. Secret Source

For a closed-source project, omit the touch command, and ensure you never create a file named git-daemon-export-ok. The repository can no longer be retrieved via the git protocol; only those with SSH access can see it. If all your repos are closed, running the git daemon is unnecessary because all communication occurs via SSH.

3.4. Bare repositories

A bare repository is so named because it has no working directory; it only contains files that are normally hidden away in the .git subdirectory. In other words, it maintains the history of a project, and never holds a snapshot of any given version.

A bare repository plays a role similar to that of the main server in a centralized version control system: the home of your project. Developers clone your project from it, and push the latest official changes to it. Typically it resides on a server that does little else but disseminate data. Development occurs in the clones, so the home repository can do without a working directory.

Many Git commands fail on bare repositories unless the GIT_DIR environment variable is set to the repository path, or the --bare option is supplied.

3.5. Push versus pull

Why did we introduce the push command, rather than rely on the familiar pull command? Firstly, pulling fails on bare repositories: instead you must *fetch*, a command we later discuss. But even if we kept a normal repository on the central server, pulling into it would still be cumbersome. We would have to login to the server first, and give the pull command the network address of the machine we're pulling from. Firewalls may interfere, and what if we have no shell access to the server in the first place?

However, apart from this case, we discourage pushing into a repository, because confusion can ensue when the destination has a working directory.

In short, while learning Git, only push when the target is a bare repository; otherwise pull.

3.6. Forking a Project

Sick of the way a project is being run? Think you could do a better job? Then on your server:

```
$ git clone git://main.server/path/to/files
```

Next, tell everyone about your fork of the project at your server.

At any later time, you can merge in the changes from the original project with:

```
$ git pull
```

3.7. Ultimate Backups

Want numerous tamper-proof geographically diverse redundant archives? If your project has many developers, don't do anything! Every clone of your code is effectively a backup. Not just of the current state, but of your project's entire history. Thanks to cryptographic hashing, if anyone's clone becomes corrupted, it will be spotted as soon as they try to communicate with others.

If your project is not so popular, find as many servers as you can to host clones.

The truly paranoid should always write down the latest 20-byte SHA1 hash of the HEAD somewhere safe. It has to be safe, not private. For example, publishing it in a newspaper would work well, because it's hard for an attacker to alter every copy of a newspaper.

3.8. Light-Speed Multitask

Say you want to work on several features in parallel. Then commit your project and run:

```
$ git clone . /some/new/directory
```

Thanks to hardlinking (http://en.wikipedia.org/wiki/Hard_link), local clones require less time and space than a plain backup.

You can now work on two independent features simultaneously. For example, you can edit one clone while the other is compiling. At any time, you can commit and pull changes from the other clone:

```
$ git pull /the/other/clone HEAD
```

3.9. Guerilla Version Control

Are you working on a project that uses some other version control system, and you sorely miss Git? Then initialize a Git repository in your working directory:

```
$ git init
```

```
$ git add .
```

```
$ git commit -m "Initial commit"
```

then clone it:

```
$ git clone . /some/new/directory
```

Now go to the new directory and work here instead, using Git to your heart's content. Once in a while, you'll want to sync with everyone else, in which case go to the original directory, sync using the other version control system, and type:

```
$ git add .
```

```
$ git commit -m "Sync with everyone else"
```

Then go to the new directory and run:

```
$ git commit -a -m "Description of my changes"
```

```
$ git pull
```

The procedure for giving your changes to everyone else depends on the other version control system. The new directory contains the files with your changes. Run whatever commands of the other version control system are needed to upload them to the central repository.

Subversion, perhaps the best centralized version control system, is used by countless projects. The git svn command automates the above for Subversion repositories, and can also be used to export a Git project to a Subversion repository

(<http://google-opensource.blogspot.com/2008/05/export-git-project-to-google-code.html>).

3.10. Mercurial

Mercurial is a similar version control system that can almost seamlessly work in tandem with Git. With the hg-git plugin, a Mercurial user can losslessly push to and pull from a Git repository.

Obtain the hg-git plugin with Git:

```
$ git clone git://github.com/schacon/hg-git.git
```

or Mercurial:

```
$ hg clone http://bitbucket.org/durin42/hg-git/
```

Sadly, I am unaware of an analogous plugin for Git. For this reason, I advocate Git over Mercurial for the main repository, even if you prefer Mercurial. With a Mercurial project, usually a volunteer maintains a

Chapter 4. Branch Wizardry

Instant branching and merging are the most lethal of Git's killer features.

Problem: External factors inevitably necessitate context switching. A severe bug manifests in the released version without warning. The deadline for a certain feature is moved closer. A developer whose help you need for a key section of the project is about to leave. In all cases, you must abruptly drop what you are doing and focus on a completely different task.

Interrupting your train of thought can be detrimental to your productivity, and the more cumbersome it is to switch contexts, the greater the loss. With centralized version control we must download a fresh working copy from the central server. Distributed systems fare better, as we can clone the desired version locally.

But cloning still entails copying the whole working directory as well as the entire history up to the given point. Even though Git reduces the cost of this with file sharing and hard links, the project files themselves must be recreated in their entirety in the new working directory.

Solution: Git has a better tool for these situations that is much faster and more space-efficient than cloning: git branch.

With this magic word, the files in your directory suddenly shapeshift from one version to another. This transformation can do more than merely go back or forward in history. Your files can morph from the last release to the experimental version to the current development version to your friend's version and so on.

4.1. The Boss Key

Ever played one of those games where at the push of a button ("the boss key"), the screen would instantly display a spreadsheet or something? So if the boss walked in the office while you were playing the game you could quickly hide it away?

In some directory:

```
$ echo "I'm smarter than my boss" > myfile.txt
```

```
$ git init
```

```
$ git add .
```

```
$ git commit -m "Initial commit"
```

We have created a Git repository that tracks one text file containing a certain message. Now type:

```
$ git checkout -b boss # nothing seems to change after this
```

```
$ echo "My boss is smarter than me" > myfile.txt
```

```
$ git commit -a -m "Another commit"
```

It looks like we've just overwritten our file and committed it. But it's an illusion. Type:

```
$ git checkout master # switch to original version of the file
```

and hey presto! The text file is restored. And if the boss decides to snoop around this directory, type:

```
$ git checkout boss # switch to version suitable for boss' eyes
```

You can switch between the two versions of the file as much as you like, and commit to each independently.

4.2. Dirty Work

Say you're working on some feature, and for some reason, you need to go back three versions and temporarily put in a few print statements to see how something works. Then:

```
$ git commit -a
```

```
$ git checkout HEAD~3
```

Now you can add ugly temporary code all over the place. You can even commit these changes. When you're done, \$ git checkout master

to return to your original work. Observe that any uncommitted changes are carried over.

What if you wanted to save the temporary changes after all? Easy:

```
$ git checkout -b dirty
```

and commit before switching back to the master branch. Whenever you want to return to the dirty changes, simply type:

```
$ git checkout dirty
```

We touched upon this command in an earlier chapter, when discussing loading old states. At last we can tell the whole story: the files change to the requested state, but we must leave the master branch. Any commits made from now on take your files down a different road, which can be named later.

In other words, after checking out an old state, Git automatically puts you in a new, unnamed branch, which can be named and saved with git checkout -b.

4.3. Quick Fixes

You're in the middle of something when you are told to drop everything and fix a newly discovered bug in commit 1b6d...:

```
$ git commit -a
```

```
$ git checkout -b fixes 1b6d
```

Then once you've fixed the bug:

```
$ git commit -a -m "Bug fixed" $ git checkout
```

```
master
```

and resume work on your original task. You can even *merge* in the freshly baked bugfix:

```
$ git merge fixes
```

4.4. Merging

With some version control systems, creating branches is easy but merging them back together is tough. With Git, merging is so trivial that you might be unaware of it happening.

We actually encountered merging long ago. The pull command in fact *fetches* commits and then merges them into your current branch. If you have no local changes, then the merge is a *fast forward*, a degenerate case akin to fetching the latest version in a centralized version control system. But if you do have local changes, Git will automatically merge, and report any conflicts.

Ordinarily, a commit has exactly one *parent commit*, namely, the previous commit. Merging branches together produces a commit with at least two parents. This begs the question: what commit does HEAD~10 really refer to? A commit could have multiple parents, so which one do we follow?

It turns out this notation chooses the first parent every time. This is desirable because the current branch becomes the first parent during a merge; frequently you're only concerned with the changes you made in the current branch, as opposed to changes merged in from other branches.

You can refer to a specific parent with a caret. For example, to show the logs from the second parent:

```
$ git log HEAD^2
```

You may omit the number for the first parent. For example, to show the differences with the first parent:

```
$ git diff HEAD^
```

You can combine this notation with other types. For example:

```
$ git checkout 1b6d^^2~10 -b ancient
```

starts a new branch "ancient" representing the state 10 commits back from the second parent of the first parent of the commit starting with 1b6d.

4.5. Uninterrupted Workflow

Often in hardware projects, the second step of a plan must await the completion of the first step. A car undergoing repairs might sit idly in a garage until a particular part arrives from the factory. A prototype might wait for a chip to be fabricated before construction can continue.

Software projects can be similar. The second part of a new feature may have to wait until the first part has been released and tested. Some projects require your code to be reviewed before accepting it, so you might wait until the first part is approved before starting the second part.

Thanks to painless branching and merging, we can bend the rules and work on Part II before Part I is officially ready. Suppose you have committed Part I and sent it for review. Let's say you're in the master branch. Then branch off:

```
$ git checkout -b part2
```

Next, work on Part II, committing your changes along the way. To err is human, and often you'll want to go back and fix something in Part I. If you're lucky, or very good, you can skip these lines.

```
$ git checkout master # Go back to Part I.
```

```
$ fix_problem
```

```
    $ git commit -a                # Commit the fixes.
```

```
    $ git checkout part2          # Go back to Part II.
```

```
    $ git merge master            # Merge in those fixes.
```

Eventually, Part I is approved:

```
$ git checkout master # Go back to Part I.
```

```
    $ submit_files                # Release to the world!
```

```
    $ git merge part2            # Merge in Part II.
```

```
$ git branch -d part2 # Delete "part2" branch.
```

Now you're in the master branch again, with Part II in the working directory.

It's easy to extend this trick for any number of parts. It's also easy to branch off retroactively: suppose you belatedly realize you should have created a branch 7 commits ago. Then type:

```
$ git branch -m master part2 # Rename "master" branch to "part2".
```

```
    $ git branch master HEAD~7    # Create new "master", 7 commits upstream.
```

The master branch now contains just Part I, and the part2 branch contains the rest. We are in the latter branch; we created master without switching to it, because we want to continue work on part2. This is unusual. Until now, we've been switching to branches immediately after creation, as in:

```
$ git checkout HEAD~7 -b master # Create a branch, and switch to it.
```

4.6. Reorganizing a Medley

Perhaps you like to work on all aspects of a project in the same branch. You want to keep works-in-progress to yourself and want others to see your commits only when they have been neatly organized. Start a couple of branches:

```
    $ git branch sanitized        # Create a branch for sanitized commits.
```

```
$ git checkout -b medley # Create and switch to a branch to work in.
```

Next, work on anything: fix bugs, add features, add temporary code, and so forth, committing often along the way. Then:

```
$ git checkout sanitized
```

```
$ git cherry-pick medley^^
```

applies the grandparent of the head commit of the "medley" branch to the "sanitized" branch. With appropriate cherry-picks you can construct a branch that contains only permanent code, and has related commits grouped together.

4.7. Managing Branches

List all branches by typing:

```
$ git branch
```

By default, you start in a branch named "master". Some advocate leaving the "master" branch untouched and creating new branches for your own edits.

The -d and -m options allow you to delete and move (rename) branches. See git help branch.

The "master" branch is a useful custom. Others may assume that your repository has a branch with this name, and that it contains the official version of your project. Although you can rename or obliterate the "master" branch, you might as well respect this convention.

4.8. Temporary Branches

After a while you may realize you are creating short-lived branches frequently for similar reasons: every other branch merely serves to save the current state so you can briefly hop back to an older state to fix a high-priority bug or something.

It's analogous to changing the TV channel temporarily to see what else is on. But instead of pushing a couple of buttons, you have to create, check out, merge, and delete temporary branches. Luckily, Git has a shortcut that is as convenient as a TV remote control:

```
$ git stash
```

This saves the current state in a temporary location (a *stash*) and restores the previous state. Your working directory appears exactly as it was before you started editing, and you can fix bugs, pull in upstream changes, and so on. When you want to go back to the stashed state, type:

```
$ git stash apply # You may need to resolve some conflicts.
```

You can have multiple stashes, and manipulate them in various ways. See `git help stash`. As you may have guessed, Git maintains branches behind the scenes to perform this magic trick.

4.9. Work How You Want

You might wonder if branches are worth the bother. After all, clones are almost as fast, and you can switch between them with `cd` instead of esoteric Git commands. Consider web browsers. Why support multiple tabs as well as multiple windows? Because allowing both accommodates a wide variety of styles. Some users like to keep only one browser window open, and use tabs for multiple webpages. Others might insist on the other extreme: multiple windows with no tabs anywhere. Others still prefer something in between. Branching is like tabs for your working directory, and cloning is like opening a new browser window. These operations are fast and local, so why not experiment to find the combination that best suits you? Git lets you work exactly how you want.

Chapter 5. Lessons of History

A consequence of Git's distributed nature is that history can be edited easily. But if you tamper with the past, take care: only rewrite that part of history which you alone possess. Just as nations forever argue over who committed what atrocity, if someone else has a clone whose version of history differs to yours, you will have trouble reconciling when your trees interact. Some developers strongly feel history should be immutable, warts and all. Others feel trees should be made presentable before they are unleashed in public. Git accommodates both viewpoints. Like cloning, branching, and merging, rewriting history is simply another power Git gives you. It is up to you to use it wisely.

5.1. I Stand Corrected

Did you just commit, but wish you had typed a different message? Then run:
`$ git commit --amend`
 to change the last message. Realized you forgot to add a file? Run `git add` to add it, and then run the above command.
 Want to include a few more edits in that last commit? Then make those edits and run:
`$ git commit --amend -a`

5.2. ... And Then Some

Suppose the previous problem is ten times worse. After a lengthy session you've made a bunch of commits. But you're not quite happy with the way they're organized, and some of those commit messages could use rewording. Then type:
`$ git rebase -i HEAD~10`
 and the last 10 commits will appear in your favourite \$EDITOR. A sample excerpt:
 pick 5c6eb73 Added repo.or.cz link
 pick a311a64 Reordered analogies in "Work How You Want"
 pick 100834f Added push target to Makefile
 Older commits precede newer commits in this list, unlike the `log` command. Here, 5c6eb73 is the oldest commit, and 100834f is the newest. Then:

- Remove commits by deleting lines. Like the `revert` command, but off the record: it will be as if the commit never existed.
- Reorder commits by reordering lines.
- Replace pick with:
- edit to mark a commit for amending.
- reword to change the log message.
- squash to merge a commit with the previous one.
- fixup to merge a commit with the previous one and discard the log message.

For example, we might replace the second pick with squash:
 pick 5c6eb73 Added repo.or.cz link
 link squash a311a64 Reordered analogies in "Work How You Want"
 pick 100834f Added push target to Makefile
 After we save and quit, Git merges a311a64 into 5c6eb73. Thus squash merges into the next commit up: think "squash up".

Git then combines their log messages and presents them for editing. The command `fixup` skips this step; the squashed log message is simply discarded. If you marked a commit with `edit`, Git returns you to the past, to the oldest such commit. You can amend the old commit as described in the previous section, and even create new commits that belong here. Once you're pleased with the "retcon", go forward in time by running: `$ git rebase --continue`
 Git replays commits until the next edit, or to the present if none remain.

You can also abandon the rebase with: `$ git`

`rebase --abort`

So commit early and commit often: you can tidy up later with `rebase`.

5.3. Local Changes Last

You're working on an active project. You make some local commits over time, and then you sync with the official tree with a merge. This cycle repeats itself a few times before you're ready to push to the central tree. But now the history in your local Git clone is a messy jumble of your changes and the official changes. You'd prefer to see all your changes in one contiguous section, and after all the official changes. This is a job for `git rebase` as described above. In many cases you can use the `--onto` flag and avoid interaction. Also see `git help rebase` for detailed examples of this amazing command. You can split commits. You can even rearrange branches of a tree. Take care: `rebase` is a powerful command. For complicated rebases, first make a backup with `git clone`.

5.4. Rewriting History

Occasionally, you need the source control equivalent of airbrushing people out of official photos, erasing them from history in a Stalinesque fashion. For example, suppose we intend to release a project, but it involves a file that should be kept private for some reason. Perhaps I left my credit card number in a text file and accidentally added it to the project. Deleting the file is insufficient, for the file can be accessed from older commits. We must remove the file from all commits:
`$ git filter-branch --tree-filter 'rm top/secret/file' HEAD`
 See `git help filter-branch`, which discusses this example and gives a faster method. In general, `filter-branch` lets you alter large sections of history with a single command. Afterwards, the `.git/refs/original` directory describes the state of affairs before the operation. Check the `filter-branch` command did what you wanted, then delete this directory if you wish to run more `filter-branch` commands. Lastly, replace clones of your project with your revised version if you want to interact with them later.

5.5. Making History

Want to migrate a project to Git? If it's managed with one of the more well-known systems, then chances are someone has already written a script to export the whole history to Git. Otherwise, look up `git fast-import`, which reads text input in a specific format to create Git history from scratch. Typically a script using this command is hastily cobbled together and run once, migrating the project in a single shot.

As an example, paste the following listing into temporary file, such as `/tmp/history`:
`commit refs/heads/master committer Alice <alice@example.com> Thu, 01 Jan 1970 00:00:00 +0000 data`
`<<EOT`
 Initial commit.
`EOT`
`M 100644 inline hello.c data <<EOT`
`#include <stdio.h>`
`int main() { printf("Hello, world!\n"); return 0;`
`} EOT`
`commit refs/heads/master committer Bob <bob@example.com> Tue, 14 Mar 2000 01:59:26 -0800 data`
`<<EOT Replace printf() with write(). EOT`
`M 100644 inline hello.c data <<EOT`
`#include <unistd.h>`
`int main() { write(1, "Hello, world!\n", 14); return 0;`
`} EOT`

Then create a Git repository from this temporary file by typing:
`$ mkdir project; cd project; git init`
`$ git fast-import --date-format=rfc2822 < /tmp/history`
 You can checkout the latest version of the project with:
`$ git checkout master`

The `git fast-export` command converts any repository to the `git fast-import` format, whose output you can study for writing exporters, and also to transport repositories in a human-readable format. Indeed, these commands can send repositories of text files over text-only channels.

5.6. Where Did It All Go Wrong?

You’ve just discovered a broken feature in your program which you know for sure was working a few months ago. Argh! Where did this bug come from? If only you had been testing the feature as you developed. It’s too late for that now. However, provided you’ve been committing often, Git can pinpoint the problem:

```
$ git bisect start
$ git bisect bad HEAD
$ git bisect good 1b6d
```

Git checks out a state halfway in between. Test the feature, and if it’s still broken:

```
$ git bisect bad
```

If not, replace “bad” with “good”. Git again transports you to a state halfway between the known good and bad versions, narrowing down the possibilities. After a few iterations, this binary search will lead you to the commit that caused the trouble. Once you’ve finished your investigation, return to your original state by typing:

```
$ git bisect reset
```

Instead of testing every change by hand, automate the search by running:

```
$ git bisect run my_script
```

Git uses the return value of the given command, typically a one-off script, to decide whether a change is good or bad: the command should exit with code 0 when good, 125 when the change should be skipped, and anything else between 1 and 127 if it is bad. A negative return value aborts the bisect.

You can do much more: the help page explains how to visualize bisects, examine or replay the bisect log, and eliminate known innocent changes for a speedier search.

5.7. Who Made It All Go Wrong?

Like many other version control systems, Git has a blame command:

Blame

Chapter 6. Multiplayer Git

Initially I used Git on a private project where I was the sole developer. Amongst the commands related to Git’s distributed nature, I needed only pull and clone so could I keep the same project in different places. Later I wanted to publish my code with Git, and include changes from contributors. I had to learn how to manage projects with multiple developers from all over the world. Fortunately, this is Git’s forte, and arguably its raison d’être.

6.1. Who Am I?

Every commit has an author name and email, which is shown by git log. By default, Git uses system settings to populate these fields. To set them explicitly, type:

```
$ git config --global user.name "John Doe"
$ git config --global user.email johndoe@example.com
```

Omit the global flag to set these options only for the current repository.

6.2. Git Over SSH, HTTP

Suppose you have SSH access to a web server, but Git is not installed. Though less efficient than its native protocol, Git can communicate over HTTP.

Download, compile and install Git in your account, and create a repository in your web directory:

```
$ GIT_DIR=proj.git git init
$ cd proj.git
$ git --bare update-server-info
$ cp hooks/post-update.sample hooks/post-update
```

For older versions of Git, the copy command fails and you should run: `$ chmod`

```
a+* hooks/post-update
```

Now you can publish your latest edits via SSH from any clone: `$ git`

```
push web.server:/path/to/proj.git master
```

and anybody can get your project with:

```
$ git clone http://web.server/proj.git
```

6.3. Git Over Anything

Want to synchronize repositories without servers, or even a network connection? Need to improvise during an emergency? We’ve seen git fast-export and git fast-import can convert repositories to a single file and back. We

could shuttle such files back and forth to transport git repositories over any medium, but a more efficient tool is git bundle.

The sender creates a *bundle*:

```
$ git bundle create somefile HEAD
```

then transports the bundle, somefile, to the other party somehow: email, thumb drive, an xxd printout and an OCR scanner, reading bits over the phone, smoke signals, etc. The receiver retrieves commits from the bundle by typing:

```
$ git pull somefile
```

The receiver can even do this from an empty repository. Despite its size, somefile contains the entire original git repository.

In larger projects, eliminate waste by bundling only changes the other repository lacks. For example, suppose the commit “1b6d...” is the most recent commit shared by both parties:

```
$ git bundle create somefile HEAD ^1b6d
```

If done frequently, one could easily forget which commit was last sent. The help page suggests using tags to solve this. Namely, after you send a bundle, type:

```
$ git tag -f lastbundle HEAD
```

and create new refresher bundles with:

```
$ git bundle create newbundle HEAD ^lastbundle
```

6.4. Patches: The Global Currency

Patches are text representations of your changes that can be easily understood by computers and humans alike.

This gives them universal appeal. You can email a patch to developers no matter what version control system they’re using. As long as your audience can read their email, they can see your edits.

Similarly, on your side, all you require is an email account: there’s no need to setup an online Git repository.

Recall from the first chapter:

```
$ git diff 1b6d > my.patch
```

outputs a patch which can be pasted into an email for discussion. In a Git repository, type:

```
$ git apply < my.patch
```

to apply the patch.

In more formal settings, when author names and perhaps signatures should be recorded, generate the corresponding patches past a certain point by typing:

```
$ git format-patch 1b6d
```

The resulting files can be given to git-send-email, or sent by hand. You can also specify a range of commits:

```
$ git format-patch 1b6d..HEAD^..
```

On the receiving end, save an email to a file, then type:

```
$ git am < email.txt
```

This applies the incoming patch and also creates a commit, including information such as the author.

With a browser email client, you may need to click a button to see the email in its raw original form before saving the patch to a file.

There are slight differences for mbox-based email clients, but if you use one of these, you’re probably the sort of person who can figure them out easily without reading tutorials!

6.5. Sorry, We’ve Moved

After cloning a repository, running git push or git pull will automatically push to or pull from the original URL. How does Git do this? The secret lies in config options created with the clone. Let’s take a peek:

```
$ git config --list
```

The remote.origin.url option controls the source URL; “origin” is a nickname given to the source repository. As with the “master” branch convention, we may change or delete this nickname but there is usually no reason for doing so.

If the original repository moves, we can update the URL via:

```
$ git config remote.origin.url git://new.url/proj.git
```

The branch.master.merge option specifies the default remote branch in a git pull. During the initial clone, it is set to the current branch of the source repository, so even if the HEAD of the source repository subsequently moves to a different branch, a later pull will faithfully follow the original branch.

This option only applies to the repository we first cloned from, which is recorded in the option branch.master.remote. If we pull in from other repositories we must explicitly state which branch we want:

```
$ git pull git://example.com/other.git master
```

The above explains why some of our earlier push and pull examples had no arguments.

6.6. Remote Branches

When you clone a repository, you also clone all its branches. You may not have noticed this because Git hides them away; you must ask for them specifically. This prevents branches in the remote repository from interfering with your branches, and also makes Git easier for beginners.

List the remote branches with:

```
$ git branch -r
```

You should see something like:

```
origin/HEAD origin/master
origin/experimental
```

These represent branches and the HEAD of the remote repository, and can be used in regular Git commands. For example, suppose you have made many commits, and wish to compare against the last fetched version. You could search through the logs for the appropriate SHA1 hash, but it's much easier to type:

```
$ git diff origin/HEAD
```

Or you can see what the "experimental" branch has been up to:

```
$ git log origin/experimental
```

6.7. Multiple Remotes

Suppose two other developers are working on our project, and we want to keep tabs on both. We can follow more than one repository at a time with:

```
$ git remote add other git://example.com/some_repo.git
```

```
$ git pull other some_branch
```

Now we have merged in a branch from the second repository, and we have easy access to all branches of all repositories:

```
$ git diff origin/experimental^ other/some_branch~5
```

But what if we just want to compare their changes without affecting our own work? In other words, we want to examine their branches without having their changes invade our working directory. Then rather than pull, run:

```
$ git fetch          # Fetch from origin, the default.
```

```
$ git fetch other # Fetch from the second programmer.
```

This just fetches histories. Although the working directory remains untouched, we can refer to any branch of any repository in a Git command because we now possess a local copy.

Recall that behind the scenes, a pull is simply a fetch then merge. Usually we pull because we want to merge the latest commit after a fetch; this situation is a notable exception.

See `git help remote` for how to remove remote repositories, ignore certain branches, and more.

Chapter 7. Git Grandmastery

By now, you should be able to navigate the git help pages and understand almost everything. However, pinpointing the exact command required to solve a given problem can be tedious. Perhaps I can save you some time: below are some recipes I have needed in the past.

7.1. Source Releases

For my projects, Git tracks exactly the files I'd like to archive and release to users. To create a tarball of the source code, I run:

```
$ git archive --format=tar --prefix=proj-1.2.3/ HEAD
```

7.2. Commit What Changed

Telling Git when you've added, deleted and renamed files is troublesome for certain projects. Instead, you can type:

```
$ git add .
```

```
$ git add -u
```

Git will look at the files in the current directory and work out the details by itself. Instead of the second add command, run `git commit -a` if you also intend to commit at this time. See `git help ignore` for how to specify files that should be ignored.

You can perform the above in a single pass with:

```
$ git ls-files -d -m -o -z | xargs -0 git update-index --add --remove
```

The `-z` and `-0` options prevent ill side-effects from filenames containing strange characters. As this command adds ignored files, you may want to use the `-x` or `-X` option.

7.3. My Commit Is Too Big!

Have you neglected to commit for too long? Been coding furiously and forgotten about source control until now? Made a series of unrelated changes, because that's your style?

No worries. Run:

```
$ git add -p
```

For each edit you made, Git will show you the hunk of code that was changed, and ask if it should be part of the next commit. Answer with "y" or "n". You have other options, such as postponing the decision; type "?" to learn more.

Once you're satisfied, type

```
$ git commit
```

to commit precisely the changes you selected (the *staged* changes). Make sure you omit the `-a` option, otherwise Git will commit all the edits.

What if you've edited many files in many places? Reviewing each change one by one becomes frustratingly mind-numbing. In this case, use `git add -i`, whose interface is less straightforward, but more flexible. With a few keystrokes, you can stage or unstage several files at a time, or review and select changes in particular files only. Alternatively, run `git commit --interactive` which automatically commits after you're done.

7.4. The Index: Git's Staging Area

So far we have avoided Git's famous *index*, but we must now confront it to explain the above. The index is a temporary staging area. Git seldom shuttles data directly between your project and its history. Rather, Git first writes data to the index, and then copies the data in the index to its final destination.

For example, `commit -a` is really a two-step process. The first step places a snapshot of the current state of every tracked file into the index. The second step permanently records the snapshot now in the index. Committing without the `-a` option only performs the second step, and only makes sense after running commands that somehow change the index, such as `git add`.

Usually we can ignore the index and pretend we are reading straight from and writing straight to the history. On this occasion, we want finer control, so we manipulate the index. We place a snapshot of some, but not all, of our changes into the index, and then permanently record this carefully rigged snapshot.

7.5. Don't Lose Your HEAD

The HEAD tag is like a cursor that normally points at the latest commit, advancing with each new commit. Some Git commands let you move it. For example:

```
$ git reset HEAD~3
```

will move the HEAD three commits back. Thus all Git commands now act as if you hadn't made those last three commits, while your files remain in the present. See the help page for some applications.

But how can you go back to the future? The past commits know nothing of the future.

If you have the SHA1 of the original HEAD then:

```
$ git reset 1b6d
```

But suppose you never took it down? Don't worry: for commands like these, Git saves the original HEAD as a tag called `ORIG_HEAD`, and you can return safe and sound with:

```
$ git reset ORIG_HEAD
```

7.6. HEAD-hunting

Perhaps `ORIG_HEAD` isn't enough. Perhaps you've just realized you made a monumental mistake and you need to go back to an ancient commit in a long-forgotten branch.

By default, Git keeps a commit for at least two weeks, even if you ordered Git to destroy the branch containing it.

The trouble is finding the appropriate hash. You could look at all the hash values in `.git/objects` and use trial and error to find the one you want. But there's a much easier way.

Git records every hash of a commit it computes in `.git/logs`. The subdirectory `refs` contains the history of all activity on all branches, while the file `HEAD` shows every hash value it has ever taken. The latter can be used to find hashes of commits on branches that have been accidentally lopped off.

The `reflog` command provides a friendly interface to these log files. Try

```
$ git reflog
```

Instead of cutting and pasting hashes from the `reflog`, try:

```
$ git checkout "@{10 minutes ago}"
```

Or checkout the 5th-last visited commit via:

```
$ git checkout "@{5}"
```

See the "Specifying Revisions" section of `git help rev-parse` for more.

You may wish to configure a longer grace period for doomed commits. For example:

```
$ git config gc.pruneexpire "30 days"
```

means a deleted commit will only be permanently lost once 30 days have passed and `git gc` is run.

You may also wish to disable automatic invocations of `git gc`:

```
$ git config gc.auto 0
```

in which case commits will only be deleted when you run `git gc` manually.

7.7. Building On Git

In true UNIX fashion, Git's design allows it to be easily used as a low-level component of other programs, such as GUI and web interfaces, alternative command-line interfaces, patch managements tools, importing and conversion tools and so on. In fact, some Git commands are themselves scripts standing on the shoulders of giants. With a little tinkering, you can customize Git to suit your preferences.

One easy trick is to use built-in Git aliases to shorten your most frequently used commands:

```
$ git config --global alias.co checkout
$ git config --global --get-regexp alias # display current aliases
alias.co checkout
$ git co foo # same as 'git checkout foo'
```

Another is to print the current branch in the prompt, or window title. Invoking

```
$ git symbolic-ref HEAD
```

shows the current branch name. In practice, you most likely want to remove the "refs/heads/" and ignore errors:

```
$ git symbolic-ref HEAD 2> /dev/null | cut -b 12-
```

The contrib subdirectory is a treasure trove of tools built on Git. In time, some of them may be promoted to official commands. On Debian and Ubuntu, this directory lives at `/usr/share/doc/git-core/contrib`.

One popular resident is `workdir/git-new-workdir`. Via clever symlinking, this script creates a new working directory whose history is shared with the original repository:

```
$ git-new-workdir an/existing/repo new/directory
```

The new directory and the files within can be thought of as a clone, except since the history is shared, the two trees automatically stay in sync. There's no need to merge, push, or pull.

7.8. Daring Stunts

These days, Git makes it difficult for the user to accidentally destroy data. But if you know what you are doing, you can override safeguards for common commands.

Checkout: Uncommitted changes cause checkout to fail. To destroy your changes, and checkout a given commit anyway, use the force flag:

```
$ git checkout -f HEAD^
```

On the other hand, if you specify particular paths for checkout, then there are no safety checks. The supplied paths are quietly overwritten. Take care if you use checkout in this manner.

Reset: Reset also fails in the presence of uncommitted changes. To force it through, run:

```
$ git reset --hard 1b6d
```

Branch: Deleting branches fails if this causes changes to be lost. To force a deletion, type:

```
$ git branch -D dead_branch # instead of -d
```

Similarly, attempting to overwrite a branch via a move fails if data loss would ensue. To force a branch move, type:

```
$ git branch -M source target # instead of -m
```

Unlike checkout and reset, these two commands defer data destruction. The changes are still stored in the `.git` subdirectory, and can be retrieved by recovering the appropriate hash from `.git/logs` (see "HEAD-hunting" above).

By default, they will be kept for at least two weeks.

Clean: Some git commands refuse to proceed because they're worried about clobbering untracked files.

If you're certain that all untracked files and directories are expendable, then delete them mercilessly with:

```
$ git clean -f -d
```

Next time, that pesky command will work!